

CSE 559A: Fall 2020

Problem Set 5

Due: Dec 15, 2020. 11:59 PM

Instructions

Please read the late submission and collaboration policy on the course website:

<http://www.cse.wustl.edu/~ayan/courses/cse559a/>

Install Anaconda for Python 3.6+ from <https://www.anaconda.com/download>. We will test all code on this distribution. You can install it locally in a separate directory without interfering with your system install of Python.

1. Complete the code files in `code/` by filling out the required functions.
2. Run each program to generate output images in `code/outputs` directory.
3. Create a PDF report in `solution.pdf` with \LaTeX by editing `solution.tex`. In particular, make sure you fill out your name/wustl key (top of the `.tex` file) to populate the page headers. Also fill out the final information section describing how long the problem set took you, who you collaborated / discussed the problem set with, as well as any online resources you used.
4. The main body of the report should contain responses to any math questions, and also include results and figures for the programming questions as asked for. These figures will often correspond to images generated by your python code in the `code/outputs/` directory.
5. Once you are done, “git add” the completed `solution.pdf` and your updated code files in `code/*.py`. Please do not add the generated output images, as these are already in your report (the git repo is setup to ignore those files). Then do a “git commit”, and a “git push”. Then, do a “git pull”, and a “git log” to verify the timestamp of your submission and the files included. These instructions are also explained in the “problem-sets” section of the course website.

As a general guideline for all problem sets: Write efficient code. While most of the points are for writing code that is correct, some points are allocated to efficiency. Above all, try to minimize the total number of multiplies / adds. For the same number of underlying operations, try to keep the use of `for` loops to a minimum (i.e., over a minimum number of indices). Instead, use convolution, element-wise operations over large arrays, calls to matrix multiply, etc.

PROBLEM 1 (Total: 25 points)

In Lecture 17, you learned about the Simple Linear Iterative Clustering (SLIC) approach to segments an image into "superpixels," or perceptually meaningful regions that still respect contours of the objects in the image. SLIC solves for the assignment of K cluster labels to N pixels ($K \ll N$) that minimizes the squared distance between an augmented (r, g, b, y, x) vector at every pixel and the mean of all augmented vectors assigned to the same label.

(a) The algorithm starts with a collection of K cluster centers initialized on an equally sampled grid across the image. Implement the `get_cluster_centers` method, which takes as an input the image and number of clusters and returns a $K \times 2$ array of cluster centers. We provide the `get_gradients` method from Problem Set 1, which you should use to refine your initial, evenly spaced, grid as explained in Lecture 17 so that none of your initial cluster centers lie on a sharp boundary. The support code will call your `get_cluster_centers` method and create the output images 'outputs/probl1_K_centers.jpg' for multiple values of K . Include these images in your report. **(10 points)**.

(b) Implement the `slic` function. It takes as input the image, number of clusters and cluster centers computed in part (a). It should return a $W \times H$ image where each pixel is assigned a label, $\{0, 1, \dots, K - 1\}$. Your implementation should not search over all possible assignments of clusters to all possible pixels—instead, you should only consider assigning a cluster label to those pixels in a $2S \times 2S$ local window around that cluster center ($S = \sqrt{N/K}$) as discussed in Lecture 17. You should also experiment with different relative weights of the spatial component in the augmented vector as discussed in the lecture.

The support code will run your `slic` function with multiple values of K . Include in your report the output images, 'outputs/probl1_K.jpg' for each different value of K . Also include the `spatial_weight` that you selected and explain why you selected that weight. **(15 points)**.

NOTE: For the purpose of this assignment, you do not need to convert the RGB images to LAB color space.

PROBLEM 2 (Total: 15 points)

We begin by looking at our own implementation of an autograd system as described in the class. Most of the framework is implemented in `edf.py`. It is then used for training a classifier to label hand-written digits as 0 – 10 (this is a smaller version of the MNIST dataset), using a neural-network with a single hidden layer in `mnist.py`. We strongly suggest you begin by looking at the existing code and trying to understand the implementation of various parts of the gradient computation and update pipeline.

(a) Try different parameters of batch size, learning rate, and number of hidden units, and comment on the relative performance of the classifier. Keep an eye on both the soft-max loss as well as the accuracy, for both the train and val sets. Look at the `xavier` function being used to initialize the weights. Why are the limits of the uniform distribution being chosen in that way? **(5 points)**.

(b) The provided code implements batched stochastic gradient descent. Extend this to use momentum. Define the `init_momentum` and `momentum` functions in `edf.py`, and then un-comment the corresponding lines in `mnist.py`. Comment on how the behavior of training changes (consider combining momentum with different batch sizes and learning rates). **(10 points)**

PROBLEM 3 (Total: 60 points)

It is especially important that for this question you use no external resources beyond standard python and numpy/scipy. Do not search for or use libraries or code-snippets for implementing convolutional layers. All the information you need to answer this question should be in the lecture slides.

Implement a 2D convolution layer, as a class `conv2` in `edf.py`. The function will take two nodes as input, corresponding to an “image” f which is an array of size $(\text{Batch Size}) \times H \times W \times (\text{Input-Channels})$, and a kernel k of size $K_1 \times K_2 \times (\text{Input-Channels}) \times (\text{Output-Channels})$. The layer should produce the output of a “valid” convolution (or rather a correlation) without padding, defined as follows:

$$g[b, y, x, c_2] = \sum_{k_y} \sum_{k_x} \sum_{c_1} f[b, y + k_y, x + k_x, c_1] k[k_y, k_x, c_1, c_2].$$

Implement the forward function (**20 points**), as well as the backward function that back-propagates gradients to both the input image (**20 points**), and to the kernel (**20 points**). Test this layer with the `mnist_conv.py` file which attempts MNIST classification with a convolution layer as its single hidden layer.

(Extra Credit +10 points): To keep things simple, we made downsampling a separate layer in `edf.py`. But this also means that the conv layer is doing many more computations than it has to. Try to write a more efficient version of the conv layer that takes “stride” as input, and produces a downsampled convolved output. To test this layer, modify `mnist_conv.py` to use your new layer with the stride option, instead of calling `conv` then `down2`. Note that to receive this extra-credit, your solution CAN NOT loop over pixels (i.e., spatial locations) in either the forward or backward pass.